

## 1 FIND - S

Implement and demonstrate the FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from a .CSV file.

```
import csv

with open('find-s-training-examples.csv') as csvfile:
    data = [line[:-1] for line in csv.reader(csvfile) if line[-1] == 'Y']
print('Positive training examples are: {}'.format(data))

S = ['$'] * len(data[0])

print('\nOutput at each step is {}'.format(S))

for example in data:
    i = 0
    for feature in example:
        S[i] = feature if S[i] == '$' or S[i] == feature else '?'
        i += 1
    print(S)
```

## 2 CANDIDATE ELIMINATION

For a given set of training data examples stored in a .CSV file, implement and demonstrate the Candidate-Elimination algorithm to output a description of the set of all hypotheses consistent with the training examples.

```
def consistent(h1, h2)
def candidateElimination()
```

```
import csv

with open('candidate-elimination-training-examples.csv') as file:
    data = [tuple(line) for line in csv.reader(file)]
```

```

D = []
for i in range(len(data[0])):
    D.append(list(set([ele[i] for ele in data])))

def consistent(h1, h2):
    for x, y in zip(h1, h2):
        if not (x == '?' or (x != '$' and (x == y or y == '$'))):
            return False
    return True

def candidateElimination():
    G = {('?',) * (len(data[0]) - 1),}
    S = ['$'] * (len(data[0]) - 1)
    num = 0
    print('G[{0}]'.format(num), G)
    print('S[{0}]'.format(num), S)
    for item in data:
        num += 1
        inp, res = item[:-1], item[-1]
        if res in 'Yy':
            G = {g for g in G if consistent(g, inp)}
            i = 0
            for s, x in zip(S, inp):
                if s != x:
                    S[i] = '?' if s != '$' else x
                i += 1
        else:
            S = S
            Gprev = G.copy()
            for g in Gprev:
                #if g not in G:
                #continue
                for i in range(len(g)):
                    if g[i] == '?':

```

```

        for val in D[i]:
            if val != inp[i] and S[i] == val:
                g_new = g[:i] + (val, ) + g[i + 1:]
                G.add(g_new)
            else:
                G.add(g)
        G.difference_update([h for h in G if any([consistent(h, g1) for g1 in
G if g1 != h])])
        print('G[{0}]'.format(num), G)
        print('S[{0}]'.format(num), S)

candidateElimination()

```

### 3 ID3

Write a program to demonstrate the working of the decision tree based ID3 algorithm. Use an appropriate data set for building the decision tree and apply this knowledge to classify a new sample.

```

def entropyOfList(aList):
def informationGain(tennis, splitAttributeName, targetAttributeName):
def id3(tennis, targetAttributeName, attributeNames, defaultClass = None):
def classify(instance, tree, default = None):

```

```

import math
from collections import Counter
from pprint import pprint
from pandas import DataFrame

tennis = DataFrame.from_csv('id3-training-examples.csv')
print('PlayTennis dataset:', tennis)

def entropyOfList(aList):
    cnt = Counter(aList)
    probs = [x / len(aList) for x in cnt.values()]
    entropy = sum([-prob * math.log(prob, 2) for prob in probs])
    return entropy

```

```
print('\nEntropy of PlayTennis dataset: %.4f' % entropyOfList(tennis['PlayTennis']))
```

```
def informationGain(tennis, splitAttributeName, targetAttributeName):
```

```
    split = tennis.groupby(splitAttributeName)
```

```
    agg_ent = split.agg({targetAttributeName: [entropyOfList, lambda x: len(x) / len(tennis)]})
```

```
    agg_ent.columns = ['Entropy', 'PropObservations']
```

```
    newEntropy = sum(agg_ent['Entropy'] * agg_ent['PropObservations'])
```

```
    oldEntropy = entropyOfList(tennis[targetAttributeName])
```

```
    return oldEntropy - newEntropy
```

```
print('\nInformation gain for Outlook: %.4f' % informationGain(tennis, 'Outlook', 'PlayTennis'))
```

```
print('Information gain for Temperature: %.4f' % informationGain(tennis, 'Temperature', 'PlayTennis'))
```

```
print('Information gain for Humidity: %.4f' % informationGain(tennis, 'Humidity', 'PlayTennis'))
```

```
print('Information gain for Wind: %.4f' % informationGain(tennis, 'Wind', 'PlayTennis'))
```

```
def id3(tennis, targetAttributeName, attributeNames, defaultClass = None):
```

```
    cnt = Counter(tennis[targetAttributeName])
```

```
    if len(cnt) == 1:
```

```
        return next(iter(cnt))
```

```
    elif tennis.empty or (not attributeNames):
```

```
        return defaultClass
```

```
    else:
```

```
        defaultClass = max(cnt.keys())
```

```
        gains = [informationGain(tennis, attr, targetAttributeName) for attr in attributeNames]
```

```
        best = attributeNames[gains.index(max(gains))]
```

```
        tree = {best: {}}
```

```
        remainingAttributeNames = [i for i in attributeNames if i != best]
```

```
        for attr, subset in tennis.groupby(best):
```

```
            tree[best][attr] = id3(subset, targetAttributeName, remainingAttributeNames, defaultClass)
```

```
return tree
```

```
attributeNames = list(tennis.columns)
print('\nList of attributes:', attributeNames)
attributeNames.remove('PlayTennis')
print('Predicting attributes:', attributeNames)
```

```
tree = id3(tennis, 'PlayTennis', attributeNames)
print('\nThe resultant decision tree is:')
pprint(tree)
```

```
attribute = next(iter(tree))
def classify(instance, tree, default = None):
    attribute = next(iter(tree))
    if instance[attribute] in tree[attribute].keys():
        result = tree[attribute][instance[attribute]]
        if isinstance(result, dict):
            return classify(instance, result)
        else:
            return result
    else:
        return default
```

```
trainData = tennis.iloc[1: -4]
testData = tennis.iloc[-4: ]
trainTree = id3(trainData, 'PlayTennis', attributeNames)
```

```
testData['predicted2'] = testData.apply(classify, axis = 1, args = (trainTree, 'Yes'))
print('\nPredicted values for sample data:\n', testData['predicted2'])
print('Accuracy: ', sum(testData['predicted2'] == testData['PlayTennis']) /
len(testData.index))
```

#### 4 BACK PROPAGATION

Build an Artificial Neural Network by implementing the Backpropagation algorithm and test the same using appropriate datasets.

```
def initializeNetwork(nInputs, nHidden, nOutputs):
def activate(weights, inputs):
def forwardPropagate(network, row):
def backwardPropagateError(network, expected):
def updateWeights(network, row, lRate):
def trainNetwork(network, dataset, lRate, nIter, nOutputs):
```

```
import math
```

```
import random
```

```
def initializeNetwork(nInputs, nHidden, nOutputs):
```

```
    network = []
```

```
    hiddenLayer = [{'weights': [random.uniform(-.5, .5) for i in range(nInputs + 1)]} for i in range(nHidden)]
```

```
    network.append(hiddenLayer)
```

```
    outputLayer = [{'weights': [random.uniform(-.5, .5) for i in range(nInputs + 1)]} for i in range(nOutputs)]
```

```
    network.append(outputLayer)
```

```
    print('The initial neural network is')
```

```
    for i, layer in zip(range(1, len(network) + 1), network):
```

```
        for j, neuron in zip(range(1, len(layer) + 1), layer):
```

```
            print('Layer[%d] Node[%d]: ' % (i, j), neuron)
```

```
    return network
```

```
def activate(weights, inputs):
```

```
    activation = weights[-1]
```

```
    for i in range(len(weights) - 1):
```

```
        activation += weights[i] * inputs[i]
```

```
    return activation
```

```
def forwardPropagate(network, row):
```

```
    inputs = row
```

```
    for layer in network:
```

```
        newInputs = []
```

```

    for neuron in layer:
        activation = activate(neuron['weights'], inputs)
        neuron['output'] = 1 / (1 + math.exp(-activation))
        newInputs.append(neuron['output'])
    inputs = newInputs
    return inputs

def backwardPropagateError(network, expected):
    for i in range(len(network) - 1, -1, -1):
        layer = network[i]
        errors = []
        if i != len(network) - 1:
            for j in range(len(layer)):
                error = 0
                for neuron in network[i + 1]:
                    error += neuron['weights'][j] * neuron['delta']
                errors.append(error)
        else:
            for j in range(len(layer)):
                errors.append(expected[j] - layer[j]['output'])
        for j in range(len(layer)):
            neuron = layer[j]
            neuron['delta'] = errors[j] * neuron['output'] * (1 -
neuron['output'])

def updateWeights(network, row, lRate):
    for i in range(len(network)):
        inputs = row[:-1]
        if i != 0:
            inputs = [neuron['output'] for neuron in network[i - 1]]
        for neuron in network[i]:
            for j in range(len(inputs)):
                neuron['weights'][j] += inputs[j] * neuron['delta'] *
lRate
            neuron['weights'][-1] += neuron['delta'] * lRate

```

```

def trainNetwork(network, dataset, lRate, nIter, nOutputs):
    for iter in range(nIter):
        sumOfErrors = 0
        for row in dataset:
            outputs = forwardPropagate(network, row)
            expected = [0 for i in range(nOutputs)]
            expected[row[-1]] = 1
            sumOfErrors += sum([(expected[i] - outputs[i]) ** 2 for i in
range(len(expected))])
            backwardPropagateError(network, expected)
            updateWeights(network, row, lRate)
        print(iter, 'Error = ', sumOfErrors)

random.seed()

dataset = [ [0, 0, 0], [0, 1, 1], [1, 0, 1], [1, 1, 1]]
#nInputs = len(dataset[0]) - 1
#nOutputs = len(set([row[-1] for row in dataset]))
network = initializeNetwork(2, 2, 2)
trainNetwork(network, dataset, 0.5, 20, 2)
print('The final neural network is')
for i, layer in zip(range(1, len(network) + 1), network):
    for j, neuron in zip(range(1, len(layer) + 1), layer):
        print('Layer[%d] Node[%d]: ' % (i, j), neuron)

```

## 5 NAÏVE BAYES

Write a program to implement the naïve Bayesian classifier for a sample training data set stored as a .CSV file. Compute the accuracy of the classifier, considering few test data sets.

```

def mean(numbers)
def stdev(numbers):
    variance = sum([pow(x - avg, 2) for x in numbers]) / float(len(numbers) - 1)
def summarize(dataset)
def calcProb(summary, item)

```



```

import csv
import math

def mean(numbers):
    return sum(numbers) / len(numbers)

def stdev(numbers):
    avg = mean(numbers)
    variance = sum([pow(x - avg, 2) for x in numbers]) / (len(numbers) - 1)
    return math.sqrt(variance)

def summarize(dataset):
    summaries = [(mean(attribute), stdev(attribute)) for attribute in zip(*dataset)]
    del summaries[-1]
    return summaries

def calcProb(summary, item):
    prob = 1
    for i in range(len(summary)):
        x = item[i]
        mean, stdev = summary[i]
        exponent = math.exp(-pow(x - mean, 2) / (2 * stdev ** 2))
        final = exponent / (math.sqrt(2 * math.pi) * stdev)
        prob *= final
    return prob

with open('naive-bayes-training-examples.csv') as file:
    data = [line for line in csv.reader(file)]
for i in range(len(data)):
    data[i] = [float(x) for x in data[i]]

split = int(0.90 * len(data))
train = data[:split]
test = data[split:]

```

```

print('\nTotal number of hypotheses:', len(data))
print('Number of hypotheses in training data:', len(train))
print('Number of hypotheses in test data:', len(test))
print("\nThe values assumed for the concept learning attributes are:")
print("OUTLOOK: Sunny = 1, Overcast = 2 and Rain = 3\nTEMPERATURE: Hot = 1, Mild = 2
and Cool = 3\nHUMIDITY: High = 1 and Normal = 2\nWIND: Weak = 1 and Strong = 2")
print("TARGET CONCEPT: PlayTennis where Yes = 10 and No = 5")

print("\nTraining dataset:")
for x in train:
    print(x)
print("\nTest dataset:")
for x in test:
    print(x)

yes = []
no = []
for i in range(len(train)):
    if data[i][-1] == 5.0:
        no.append(data[i])
    else:
        yes.append(data[i])

yes = summarize(yes)
no = summarize(no)

predictions = []
for item in test:
    yesProb = calcProb(yes, item)
    noProb = calcProb(no, item)
    predictions.append(10.0 if(yesProb > noProb) else 5.0)

correct = 0
for i in range(len(test)):

```

```

    if(test[i][-1] == predictions[i]):
        correct += 1

print("\nActual values are:")
for i in range(len(test)):
    print(test[i][-1], end=" ")
print("\nPredicted values are:")
for i in range(len(predictions)):
    print(predictions[i], end=" ")
print("\nAccuracy is %.1f%%" % ((correct / len(test)) * 100))

```

## 6 NAÏVE BAYES TEXT CLASSIFIER

Assuming a set of documents that need to be classified, use the naïve Bayesian Classifier model to perform this task. Built-in Java classes/API can be used to write the program. Calculate the accuracy, precision, and recall for your data set.

```

from sklearn.datasets import fetch_20newsgroups
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
from sklearn.naive_bayes import MultinomialNB

train = fetch_20newsgroups(subset = 'train', shuffle = True)
print('The categories of 20NewsGroups are:')
for cat in train.target_names:
    print(cat)

categories = ['alt.atheism', 'soc.religion.christian', 'comp.graphics', 'sci.med']
train = fetch_20newsgroups(subset = 'train', categories = categories, shuffle = True)
test = fetch_20newsgroups(subset = 'test', categories = categories, shuffle = True)

countVectorizer = CountVectorizer()
traintf = countVectorizer.fit_transform(train.data)
print('\ntf train count:', traintf.shape)
testtf = countVectorizer.transform(test.data)

```

```

print('tf test count:', testtf.shape)

tfidftransformer = TfidfTransformer()
traintfidf = tfidftransformer.fit_transform(traintf)
print('\ntf train count:', traintf.shape)
testtfidf = tfidftransformer.transform(testtf)
print('tf test count:', testtf.shape)

model = MultinomialNB()
model.fit(traintfidf, train.target)
predicted = model.predict(testtfidf)

print('Accuracy score:', accuracy_score(test.target, predicted))
print(classification_report(test.target, predicted, target_names = test.target_names))
print('Confusion Matrix: ', confusion_matrix(test.target, predicted))

```

## 7 BAYESIAN BELIEF NETWORK

Write a program to construct a Bayesian network considering medical data. Use this model to demonstrate the diagnosis of heart patients using standard Heart Disease Data Set. You can use Java/Python ML library classes/API.

```

import numpy as np
import pandas as pd
import urllib
from urllib.request import urlopen
import pgmpy
from pgmpy.models import BayesianModel
from pgmpy.estimators import MaximumLikelihoodEstimator
from pgmpy.inference import VariableElimination

url = 'http://archive.ics.uci.edu/ml/machine-learning-databases/heart-disease/processed.hungarian.data'
np.set_printoptions(threshold = np.nan)

```

```
names = ['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach', 'exang',
'oldpeak', 'slope', 'ca', 'thal', 'heartdisease']
```

```
heartDisease = pd.read_csv(urlopen(url), names = names)
print(heartDisease.head())
```

```
del heartDisease['oldpeak']
del heartDisease['slope']
del heartDisease['ca']
del heartDisease['thal']
```

```
heartDisease = heartDisease.replace('?', np.nan)
print(heartDisease.dtypes)
```

```
model = BayesianModel([('age', 'trestbps'), ('age', 'fbs'), ('sex', 'trestbps'),
('sex', 'trestbps'), ('exang', 'trestbps'), ('trestbps', 'heartdisease'), ('fbs',
'heartdisease'), ('heartdisease', 'restecg'), ('heartdisease', 'thalach'),
('heartdisease', 'chol')])
```

```
model.fit(heartDisease, estimator = MaximumLikelihoodEstimator)
```

```
print(model.get_cpds('age'))
print(model.get_cpds('chol'))
print(model.get_cpds('sex'))
```

```
model.get_independencies()
inference = VariableElimination(model)
```

```
q = inference.query(variables = ['heartdisease'], evidence = {'age': 28})
print(q['heartdisease'])
```

```
q = inference.query(variables = ['heartdisease'], evidence = {'chol': 100})
print(q['heartdisease'])
```

## 8 KMEANS AND EM

Apply EM algorithm to cluster a set of data stored in a .CSV file. Use the same data set for clustering using *k*-Means algorithm. Compare the results of these two algorithms and comment on the quality of clustering. You can add Java/Python ML library classes/API in the program.

```
import numpy as np
import pandas as pd
from sklearn.cluster import KMeans
from sklearn.mixture import GaussianMixture
from sklearn.datasets import load_iris
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn import preprocessing
import matplotlib.pyplot as plt

l1 = [0, 1, 2]
def rename(S):
    l2 = []
    for i in S:
        if i not in l2:
            l2.append(i)
    for i in S:
        pos = l2.index(i)
        i = l1[pos]
    return S

iris = load_iris()
print('Data', iris.data)
print('Target names:', iris.target_names)
print('Target:', iris.target)

x = pd.DataFrame(iris.data)
x.columns = ['SepalLength', 'SepalWidth', 'PetalLength', 'PetalWidth']
y = pd.DataFrame(iris.target)
y.columns = ['Targets']
```

```

model = KMeans(n_clusters = 3)
model.fit(x)

plt.figure(figsize = (14, 7))
colormap = np.array(['red', 'lime', 'black'])
plt.subplot(1, 2, 1)
plt.scatter(x.PetalLength, x.PetalWidth, c = colormap[y.Targets], s = 40)
plt.title('Real Classification')
plt.subplot(1, 2, 2)
plt.scatter(x.PetalLength, x.PetalWidth, c = colormap[model.labels_], s = 40)
plt.title('KMeans Classification')
plt.show()

km = rename(model.labels_)
print('What KMeans thought:', km)
print('Accuracy score of KMeans:', accuracy_score(y, km))
print('Confusion matrix of KMeans:', confusion_matrix(y, km))

scaler = preprocessing.StandardScaler()
scaler.fit(x)
xsa = scaler.transform(x)
xs = pd.DataFrame(xsa, columns = x.columns)
print('\n', xs.sample(5))

gmm = GaussianMixture(n_components = 3)
gmm.fit(xs)
y_cluster_gmm = gmm.predict(xs)

plt.subplot(1, 2, 1)
plt.scatter(x.PetalLength, x.PetalWidth, c = colormap[y_cluster_gmm], s = 40)
plt.title('GMM Classification')
plt.show()

em = rename(y_cluster_gmm)

```

```
print('What EM thought:', km)
print('Accuracy score of EM:', accuracy_score(y, em))
print('Confusion matrix of EM:', confusion_matrix(y, em))
```

## 9 K NEAREST NEIGHBOUR

Write a program to implement *k*-Nearest Neighbour algorithm to classify the iris data set. Print both correct and wrong predictions. Java/Python ML library classes can be used for this problem.

```
import numpy as np
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier

dataset = load_iris()
print('IRIS FEATURES | TARGET NAMES:', dataset.target_names)
print('\nData:', dataset["data"])
print('\nTarget', dataset["target"])

xtrain, xtest, ytrain, ytest = train_test_split(dataset["data"], dataset["target"],
random_state = 0)

print("\nX TRAIN \n", xtrain)
print("\nX TEST \n", xtest)
print("\nY TRAIN \n", ytrain)
print("\nY TEST \n", ytest)

kn = KNeighborsClassifier(n_neighbors = 1)
kn.fit(xtrain, ytrain)

predictions = kn.predict(xtest)
for i in range(len(xtest)):
    print("\nActual: {0} {1} \nPredicted: {2} {3}".format(ytest[i],
dataset["target_names"][ytest[i]], predictions, dataset["target_names"][predictions]))
print("\nTEST SCORE[ACCURACY]: {:.2f}\n".format(kn.score(xtest, ytest)))
```



**10 REGRESSION - BOKEH IS LIFE!!!**

Implement the non-parametric Locally Weighted Regression algorithm in order to fit data points. Select appropriate data set for your experiment and draw graphs.

```
import numpy as np
from bokeh.layouts import gridplot
from bokeh.plotting import figure, show

def local_regression(x0, X, Y, tau):
    x0 = np.r_[1, x0]
    X = np.c_[np.ones(len(X)), X]
    xw = X.T * np.exp(np.sum((X - x0) ** 2, axis = 1) / (-2 * tau ** 2))
    beta = np.linalg.pinv(xw @ X) @ xw @ Y
    return x0 @ beta

n = 1000
X = np.linspace(-3, 3, num = n)
print('The dataset(10 samples) X is:', X[1: 10])

Y = np.log(np.abs(X ** 2 - 1) + 0.5)
print('\nThe fitted curve dataset(10 samples) Y is:', Y[1: 10])

X += np.random.normal(scale = 0.1, size = n)
print('\nThe normalized dataset(10 samples) X is:', X[1: 10])

domain = np.linspace(-3, 3, num = 300)
print('\nThe domain (10 samples) is:', domain[1: 10])

def plot_lwr(tau):
    prediction = [local_regression(x0, X, Y, tau) for x0 in domain]
    plot = figure(plot_width = 400, plot_height = 400)
    plot.title.text = 'tau: %g' % tau
    plot.scatter(X, Y, alpha = 0.3)
```

```
plot.line(domain, prediction, line_width = 2, color = 'red')  
return plot
```

```
show(gridplot([[plot_lwr(10), plot_lwr(1)], [plot_lwr(.1), plot_lwr(.01)]]))
```